



UNIVERSITY OF  
CENTRAL  
MISSOURI®

DSA 5600 – NoSQL Database Systems  
Section 3: Apache Cassandra

Instructor: Yutong Zhao



# What is Cassandra?

- **Apache Cassandra** is a **distributed, NoSQL database** designed for handling large volumes of structured data across many servers without a single point of failure.
- Originally developed at **Facebook** to power their inbox search, later open-sourced and adopted by the **Apache Software Foundation**.
- **Key Characteristics:**
  - **Scalability:** Easily scales horizontally by adding more nodes.
  - **High Availability:** No single point of failure; data is replicated across multiple nodes.
  - **NoSQL Model:** Uses a schema-less, column-family-based storage model.
  - **Optimized for Write-Intensive Workloads:** Handles high-speed inserts and updates efficiently.
  - **Eventual Consistency:** Ensures availability over strict consistency (CAP theorem).
- **Used By:** Netflix, Twitter, eBay, Uber, and many more for high-performance, globally distributed applications.



# Key Features of Cassandra (1)

## 1. Distributed & Decentralized Architecture

- No master-slave structure; all nodes are equal (peer-to-peer model).
- Ensures high availability and fault tolerance.

## 2. Horizontal Scalability

- Easily add more nodes to handle increasing data loads.
- Supports massive datasets with linear scalability.

## 3. High Availability & Fault Tolerance

- Data is replicated across multiple nodes.
- If a node fails, another node takes over without downtime.

## 4. Tunable Consistency

- Supports different levels of consistency ( `ONE` , `QUORUM` , `ALL` ).
- Balances availability vs. consistency based on application needs.



# Key Features of Cassandra (2)

## 5. Optimized for High Write Performance

- Designed for fast inserts and updates with **log-structured storage**.
- Ideal for real-time big data applications.

## 6. Flexible Schema (NoSQL Model)

- Uses a column-family-based data model.
- Allows dynamic addition of columns without altering the schema.

## 7. Built-in Data Replication

- Automatic replication across multiple data centers.
- Ensures disaster recovery and global accessibility.

## 8. Support for Multi-Datacenter Deployments

- Allows replication across geographically distributed data centers.
- Improves performance for globally distributed users.



# When to use Cassandra?

- ✓ **High Availability & Fault Tolerance** – No single point of failure, automatic failover.
- ✓ **Massive Scalability** – Handles large datasets and high transaction volumes.
- ✓ **Fast Write Performance** – Optimized for real-time applications with heavy writes.
- ✓ **Global Distribution** – Replicates data across multiple regions/data centers.
- ✓ **Schema Flexibility** – No rigid structure; adaptable to dynamic data models.



# Cassandra vs. Relational Database (RDBMS)

Feature	Cassandra (NoSQL)	Relational DB (SQL)
Architecture	Distributed, decentralized	Centralized or master-slave
Scalability	Horizontal (add more nodes)	Vertical (upgrade server)
Performance	Optimized for fast writes	Optimized for structured transactions
Schema	Flexible, schema-less	Fixed schema, predefined tables
Consistency	Eventual consistency (CAP theorem)	Strong consistency (ACID transactions)
Use Case	Large-scale, real-time applications	Traditional transactional applications

- ✓ Choose RDBMS when you need strict consistency (e.g., banking, inventory).
- ✓ Choose Cassandra when you need distributed, highly available systems (e.g., real-time analytics, IoT, messaging apps).



# Cassandra vs. NoSQL Database

Feature	Cassandra	MongoDB	HBase
Data Model	Wide-column store	Document store	Wide-column store
Scalability	Highly scalable	Moderate scalability	Good for large-scale batch processing
Write Performance	High	Moderate	High
Read Performance	Moderate	High (indexed JSON)	High (batch reads)
Best Use Case	High-write workloads	Document-based storage	Large-scale batch analytics

- ✓ Choose MongoDB if you need flexible JSON document storage.
- ✓ Choose HBase for big data batch processing (Hadoop integration).
- ✓ Choose Cassandra for high-throughput, real-time applications with global scalability.



# Architecture Overview: Nodes, Partitions, Replication

## **Cassandra's Distributed Architecture**

- ✓ **Peer-to-Peer Model** – No master-slave setup; all nodes are equal.
- ✓ **High Availability & Fault Tolerance** – Data is automatically distributed across nodes.
- ✓ **Ring-Based Architecture** – Nodes are logically arranged in a circle to balance load.

## **Nodes & Partitions: How Cassandra Stores Data**

- ◆ **Node** – A single machine in the Cassandra cluster.
- ◆ **Partitioning** – Data is divided into **partitions** using a **Partition Key**, ensuring even distribution.
- ◆ **Data Centers & Clusters** – Nodes are grouped into **data centers** for geographical redundancy.





# Architecture Overview: Nodes, Partitions, Replication

## **Replication: Ensuring Reliability**

- ◆ Replication Factor (RF) – Defines how many copies of data exist in the system.
- ◆ RF = 3 → Each piece of data is stored on **3 different nodes** for redundancy.
- ◆ Replication Strategies:
  - SimpleStrategy – Used for single data centers.
  - NetworkTopologyStrategy – Used for multi-region data replication.

## **The CAP Theorem & Cassandra's Trade-offs**

- ◆ Cassandra prioritizes:
  - ✓ **Availability** – Always accessible, even if some nodes fail.
  - ✓ **Partition Tolerance** – Can handle network failures without downtime.
  - ✗ **Consistency (Eventual Consistency)** – Data updates propagate across nodes over time.
- ◆ Supports **Tunable Consistency** – Choose between stronger or weaker consistency based on needs.



# Cassandra's Data Model

## ◆ How Cassandra's Data Model is Different from SQL

- No strict schemas, foreign keys, or JOIN operations like in relational databases.
- Optimized for scalability and fast queries, rather than complex relationships.
- Uses denormalization to store data based on query patterns instead of strict normalization.

## ◆ Key Components of Cassandra's Data Model

- ✓ Keyspaces – The highest-level container, similar to a database in SQL.
- ✓ Tables (Column Families) – Stores data but does not enforce a strict schema.
- ✓ Primary Keys – Unique identifier that determines data distribution.



# Key Features of Cassandra

## ◆ **Partition Keys: Data Distribution in Cassandra**

- The **Partition Key** decides which node stores the data.
- Helps distribute data **evenly** across the cluster.
- Ensures **fast and scalable queries** by keeping related data together.

## ◆ **Clustering Columns: Organizing Data Within Partitions**

- Defines how data is sorted inside a partition.
- Enables efficient **ordering and filtering** within a partition.

## ◆ **Denormalization: Why Cassandra Avoids Joins**

- Data is **duplicated** across tables to avoid costly **JOIN** operations.
- Ensures **faster queries** by storing **pre-aggregated** data.



# Creating a Keyspace

```
CREATE KEYSPACE store  
WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 3};
```

- Keyspace `store` created with SimpleStrategy replication.
- Replication Factor = 3, meaning data is stored on three different nodes for redundancy.



# Creating a Table with Primary and Clustering Keys

```
CREATE TABLE store.products (  
  product_id UUID PRIMARY KEY,  
  name TEXT,  
  price DECIMAL,  
  stock INT  
);
```

- *product\_id* is the Partition Key, which decides which node stores the data.
- The table is schema-flexible, allowing new columns to be added dynamically.



# Inserting and Querying Data

```
INSERT INTO store.products (product_id, name, price, stock)  
VALUES (uuid(), 'Laptop', 1200, 50);
```

```
SELECT * FROM store.products;
```



# CRUD Operations in Cassandra

- ✓ **Create** – Insert new data into a table.
- ✓ **Read** – Retrieve data using optimized queries.
- ✓ **Update** – Modify existing records.
- ✓ **Delete** – Remove data while maintaining system efficiency.



# Query Data with CQL

- Basic Querying with SELECT:

*-- Retrieve all records*

```
SELECT * FROM store.products;
```

*-- Retrieve specific columns*

```
SELECT name, price FROM store.products;
```

*-- Retrieve a specific product by Primary Key*

```
SELECT * FROM store.products WHERE product_id = <some-uuid>;
```





# Query Data with CQL

- Filtering Data:

```
-- Query using Partition Key (Recommended)
SELECT * FROM store.products WHERE product_id = <some-uuid>;

-- Querying with Clustering Columns
SELECT * FROM store.orders WHERE customer_id = 123
ORDER BY order_date DESC;

-- Using ALLOW FILTERING (Not Recommended for Large Datasets)
SELECT * FROM store.products WHERE price > 500 ALLOW FILTERING;
```



# Query Data with CQL

- **Aggregation & Counting Records:**

```
-- Count total records (Not Optimized)
```

```
SELECT COUNT(*) FROM store.products;
```

```
-- Count records within a partition (Recommended)
```

```
SELECT COUNT(*) FROM store.products WHERE category = 'Electronics';
```



# Materialized View

```
CREATE MATERIALIZED VIEW store.product_by_name AS  
SELECT name, product_id, price, stock  
FROM store.products  
WHERE name IS NOT NULL  
PRIMARY KEY (name, product_id);
```



# Secondary Index

```
-- Create a Secondary Index on the 'name' column  
CREATE INDEX ON store.products (name);  
  
-- Query products using the indexed column  
SELECT * FROM store.products WHERE name = 'Laptop';  
  
-- Drop the Secondary Index if no longer needed  
DROP INDEX store.products_name_idx;
```



# Performance Optimization & Data Replication

- ✓ **Replication Strategies** – Distributes data across nodes for **fault tolerance & availability**.
- ✓ **Consistency Levels** – Controls **how strongly synchronized** data needs to be across nodes.
- ✓ **Read & Write Optimization** – Techniques like **batch queries, indexing, and partitioning** for better performance.
- ✓ **Caching & Compaction** – Improves **query speed and storage efficiency** by managing memory and disk usage.



# Replication Strategies in Cassandra

```
-- SimpleStrategy (For single data centers)
CREATE KEYSPACE ecommerce
WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 3};

-- NetworkTopologyStrategy (For multi-data center setups)
CREATE KEYSPACE ecommerce
WITH replication = {'class': 'NetworkTopologyStrategy', 'DC1': 3, 'DC2': 2};
```



# Understanding Consistency Levels

Level	Reads Required	Writes Required	Best For
ONE	1 Node	1 Node	Fastest, lower consistency
QUORUM	Majority of Nodes	Majority of Nodes	Balance of availability & consistency
ALL	All Nodes	All Nodes	Strongest consistency, slowest

```
-- Writing with a specific consistency level
INSERT INTO store.products (product_id, name, price, stock)
VALUES (uuid(), 'Smartphone', 800, 150)
USING CONSISTENCY QUORUM;

-- Reading with a specific consistency level
SELECT * FROM store.products USING CONSISTENCY ONE;
```



# Tuning Read & Write Performance

*-- Creating an Index for Faster Reads (Use with caution)*

```
CREATE INDEX ON store.products (category);
```

*-- Using Materialized Views for Optimized Queries*

```
CREATE MATERIALIZED VIEW store.products_by_category AS
```

```
SELECT category, name, price
```

```
FROM store.products
```

```
WHERE category IS NOT NULL
```

```
PRIMARY KEY (category, name);
```

*-- Batched Writes for Efficiency*

```
BEGIN BATCH
```

```
    INSERT INTO store.products (product_id, name, price, stock)
```

```
    VALUES (uuid(), 'Tablet', 300, 100);
```

```
APPLY BATCH;
```





# Using Caching & Compaction Strategies

```
ALTER TABLE store.products  
WITH caching = {'keys': 'ALL', 'rows_per_partition': '10'};
```

```
ALTER TABLE store.products  
WITH compaction = {'class': 'SizeTieredCompactionStrategy'};
```



# Backup and Restore in Cassandra

## Backup Strategies

- ◆ Snapshots 📷 – Captures a full backup of SSTables (nodetool snapshot)
- ◆ Incremental Backups ⌚ – Saves only changes since the last snapshot
- ◆ Commit Log Archiving 📁 – Ensures recovery of unflushed writes
- ◆ Exporting Data 🌐 – Use `cqlsh COPY` or `nodetool flush` to create backups

## Restore Methods




- ✓ Restore from Snapshots 📁 – Copy SSTable files back and refresh the database
- ✓ Replay Commit Logs 🔄 – Recover unflushed writes after failure
- ✓ Use Sstableloader 🚀 – Import backed-up data into a new cluster

💡 *Tip: Automate backups with scripts & cloud storage for better disaster recovery!*









# Security Features in Cassandra

## Authentication & Authorization

- ◆ User Authentication  – Verify identities before accessing Cassandra
- ◆ Role-Based Access Control (RBAC)  – Assign roles & permissions
- ◆ LDAP & Kerberos Support  – Integrate with enterprise security





## Data Encryption

-  Client-to-Node Encryption  – Protects data in transit
-  Node-to-Node Encryption  – Secures inter-node communication
-  At-Rest Encryption  – Safeguards stored data using secure keys







# Monitoring and Troubleshooting Cassandra




## Monitoring Cassandra Performance

- ◆ `nodetool status`  – Check cluster health & node availability
- ◆ `nodetool tpstats`  – View active thread pools & latency stats
- ◆ `nodetool cfstats`  – Get per-table statistics (reads, writes, compaction)
- ◆ JMX & Prometheus  – Collect metrics for real-time monitoring

## Common Issues & Troubleshooting

- High Latency  – Check read/write consistency levels & tuning
- Node Failure  – Verify logs & use `nodetool repair` for recovery
- Compaction Overhead  – Adjust compaction strategy for efficiency
- Out of Memory (OOM)  – Tune heap size & GC settings

## Tools for Troubleshooting

- ✓ System Logs  – `/var/log/cassandra` for error tracking
- ✓ `nodetool describcluster`  – Check cluster-wide settings
- ✓ Tracing Queries  – Use `CONSISTENCY TRACE` for slow queries





# Advanced Data Modeling in Cassandra

## ▢ Key Principles of Data Modeling

- ◆ Denormalization Over Joins 📌 – Store redundant data to avoid costly joins
- ◆ Query-Driven Design 🎯 – Model data based on how it will be queried
- ◆ Partitioning Strategy ⚡ – Choose efficient partition keys for even data distribution
- ◆ Clustering Keys 🔗 – Define sort order for query performance




## 🔧 Best Practices for Schema Design

- ✓ Use Composite Primary Keys 🔑 – Combine partition & clustering keys wisely
- ✓ Avoid Large Partitions 🚨 – Distribute data evenly to prevent hotspots
- ✓ Leverage Indexing 📖 – Use secondary indexes or materialized views carefully
- ✓ TTL & Expiration ⌚ – Set Time-to-Live (TTL) for temporary data







# Integration with Other Tools

## Why Integrate Cassandra with Other Tools?

- ◆ Real-Time Analytics  – Process and analyze large-scale data efficiently
- ◆ Streaming Data Processing  – Handle high-velocity event streams
- ◆ Search & Indexing  – Enhance querying capabilities beyond CQL





## Common Integrations

- ◆ Apache Spark  – Distributed data processing & analytics
- ◆ Apache Kafka  – High-throughput message streaming
- ◆ Elasticsearch  – Full-text search & indexing
- ◆ Apache Flink  – Real-time event processing



# Case Studies of Cassandra in Production

## How Leading Companies Use Cassandra

- ◆ Netflix  – Handles billions of daily video streaming requests
- ◆ Instagram  – Stores and serves petabytes of user-generated content
- ◆ Uber  – Manages real-time geolocation and ride-matching
- ◆ eBay  – Powers distributed transactional data for marketplace operations



